



SPECIFICATION · V1.0

The Fact Sheet

The sixteen questions we ask an agency, and the schema those answers become — the ground truth every accuracy finding in an AI-visibility audit is measured against.

Prepared

5 August 2026

Status

Question set implemented
Schema: 1 item outstanding

Source of truth

`src/audit/factsheet/intake/questions.py`
`src/audit/factsheet/models.py`

What this is, and why it pays

A fact sheet is the reference an LLM judge measures AI answers against. It is what turns “are they mentioned?” into “here is what the AI gets wrong about them.” Mention, prominence and framing are scored straight from the answer and need no sheet at all — this document exists purely so the judge can tell true from false.

The asymmetry that governs every rule below. A wrong line in a fact sheet does not produce a missing finding. It produces a **false accusation in a document we send a stranger**. Hence: checkable facts only, blank is safe, and a schema that refuses malformed claims rather than repairing them.

RULE 01

Only falsifiable facts

A price, a partner badge, a minimum spend. Never “the leading agency.” The judge can only check things that are true or false; a marketing line is a line that can never fire.

RULE 02

Blank is safe, and it is the default

A skipped question produces zero claims, so a silent dimension can never produce a false flag. Fourteen confirmed lines beat forty with six guesses. All sixteen are skippable.

RULE 03

Negatives are where the value is

We don't build websites. No public email. Not a Premier Partner. An omission is never a flag — only a contradiction is. Seven of the sixteen lead with the negative.

RULE 04

Show the sentence they'll be quoted on

Before an answer commits, the interface renders the assertion it becomes — “*There is no free audit.*” — not the raw input, “No.”

Why confirmation, not extraction, is the whole point

Each claim carries a verification tier, and a sheet is only as strong as its **weakest** claim. A sheet assembled purely from a crawl is permanently capped at low and medium — so the findings that actually sell (*a wrong minimum spend, a partner tier you don't hold, a service you dropped*) stay switched off until a human confirms the lines they rest on. These sixteen questions are the only mechanism that reaches `client_confirmed`.

VERIFICATION TIER	MEANING	SEVERITIES IT MAY SEND
public_source_only	One public source, quoted	Low · Medium
cross_confirmed	Two independent sources agree	Low · Medium · High · Critical
client_confirmed	The client said so, on the record	Low · Medium · High · Critical

High and Critical are suppressed, not downgraded — a serious claim wearing a softer label is worse than silence.

The sixteen questions

One set, every client. The judge checks *dimensions* — identity, what you offer, what it costs, how to reach you, what you can prove — not business types, so the same sixteen questions serve an agency, its clients, and everything on its roster. Only the *for instance* line changes; never the question, never the keys, never whether a card appears. Ceiling sixteen cards; median eleven once a crawl pre-fills the confirmable ones.

The six categories

1	Who you are The name, what the business actually does, the other ways people write it, and who they mix you up with.	4 questions identity competitor_confusion
2	What you offer The list a customer would ask about — and the things you <i>don't</i> do, and what has changed lately.	3 questions missing_or_invented_feature stale
3	What it costs Prices with a basis attached, and the mandatory extras nobody puts on the pricing page.	2 questions wrong_pricing
4	How to reach you Every channel, where you can deliver, when a person is actually reachable — and the numbers and addresses still online that are no longer yours.	3 questions wrong_contact wrong_service_area wrong_hours
5	What you can prove Licences, certifications, compliance, partner tiers — and who the buyer would look at instead.	2 questions licensing
6	What AI gets wrong What you have already caught a model saying about you. The only category that is not about the business.	2 questions feeds every flag

One to five are what is true. Six is what is said. The first five are facts about the business, and a crawl can guess at most of them. The sixth is about the models, and nothing but a conversation produces it.

*Within categories two, three, four and five the **negative** half carries most of the value — we don't build websites, there is a ten-thousand minimum, no public email, not a Premier Partner. An omission is never a flag; only a contradiction is.*

Who you are identity

01 Q-WHAT-01 BATCH CONFIRM

Quick check — tap anything that's wrong.

Four facts in one tap, and the category is what an AI gets asked for.

Only what we actually found is shown. Anything blank, we just don't claim.

FOR INSTANCE — A MARKETING AGENCY

“Performance marketing agency” — not “digital services company.” The words you want an AI to use.

SECTION identity KEYS identity_name identity_website identity_founded
identity_category

02 Q-WHAT-02 LONG TEXT

In one sentence, what does {business} actually do?

This is the line every answer gets measured against.

Plain and factual. Skip the sales language — an AI can't be wrong about “the best”, only about what you do.

FOR INSTANCE — A MARKETING AGENCY

“A performance marketing agency that manages paid search, paid social, Amazon ads and ad creative for growth-stage brands.”

SECTION identity KEYS identity_what

03 Q-WHAT-03 CHIP LIST

Any other ways people write your name?

So we count a mention that spells you differently.

Misspellings, your legal name, an old name. We use these to catch a mention we'd otherwise miss — they aren't fact-checked.

FOR INSTANCE — A MARKETING AGENCY

The LLC on your invoices, a pre-merger name, an old trading name, the one-word spelling.

SECTION — KEYS identity_aliases

04 Q-WHAT-04 CHIP LIST NEGATIVE-FIRST

Is there anyone people mix you up with?

One of the things AI gets wrong most often, and one of the easiest to catch.

Same-name businesses, a former partner, a company you were spun out of.

FOR INSTANCE — A MARKETING AGENCY

A holding group you joined, a same-name agency in another city, the AI vendor behind your stack.

SECTION identity KEYS identity_not

What you provide services_pricing · features

05 Q-OFFER-01 CHIP LIST

What do you actually offer?

The list an AI is supposed to be reading off.

Name the level a customer would ask about — the categories, not every individual item. Ten or so is plenty.

FOR INSTANCE — A MARKETING AGENCY

Paid search · paid social · performance creative · Amazon ads · SEO · AEO/GEO · HubSpot · analytics setup.
Categories, not the nine SEO sub-types.

SECTION services_pricing KEYS services_offered

06 Q-OFFER-02 CHIP LIST NEGATIVE-FIRST

What do people ask for that you don't do?

Without it, nobody can catch an AI volunteering you for work you don't take.

The thing you turn down every week.

FOR INSTANCE — A MARKETING AGENCY

"We don't build websites." "We don't do brand identity." "We don't manage organic social." "No PR." The adjacent service a model will assume you also do.

SECTION services_pricing KEYS services_excluded

07 Q-OFFER-03 REPEATING PAIRS

What's new, and what have you stopped doing?

The fastest way to catch an AI still describing you as you were last year.

AI training data lags by months. The things you've stopped doing matter most — an AI recommending a service you dropped sends a customer to a dead end.

The newest thing you've launched	Text
What you've added	Repeating pairs
What you've stopped doing	Repeating pairs

FOR INSTANCE — A MARKETING AGENCY

“Added AEO/GEO, March.” “Joined the holding group, June.” “Stopped taking one-off project work.”

SECTION features KEYS features_current features_added features_removed

What it costs services_pricing

08 Q-COST-01 PRICED ROWS

What does it cost?

Prices are the single most-hallucinated thing about any business.

Exact numbers, today's. “Free”, “From \$X”, “Varies by scope” and “Quote only” are all real answers.

FOR INSTANCE — A MARKETING AGENCY

“Media management · fixed monthly retainer · per month · scaled to ad spend.” The basis is checkable even when the number is not published.

SECTION services_pricing KEYS pricing_rows

09 Q-COST-02 CHOICE NEGATIVE-FIRST

Is there anything people have to pay on top of that?

AI quotes the headline price and misses this constantly.

And: is there a free option — a free tier, a free trial, a free consult, free shipping? “No” is a valuable answer to both.

Anything paid on top	Choice	No, that's the full price / Yes —
A free option — free tier, trial, consult, shipping	Choice	No / Yes —

FOR INSTANCE — A MARKETING AGENCY

A minimum ad spend, an onboarding fee, tooling billed separately. And: is the intro audit actually free?

SECTION services_pricing KEYS pricing_mandatory_extra pricing_free_option

How to reach you contact · coverage · availability

10 Q-REACH-01 BATCH CONFIRM NEGATIVE-FIRST

This is the most important card here — tap anything that's wrong.

A wrong way to reach you is a customer you never hear about.

“We don't have a phone line” and “no public address” are real answers here, and good ones. Then: any old number, address or link still floating around?

Phone	Text
Email	Text
Booking or support link	Text
Address	Text
We don't have one of these phone · email · booking · address	Chip list
Old numbers, addresses or links still online	Chip list

FOR INSTANCE — A MARKETING AGENCY

Phone and a booking link. “There is no public email address” is a real answer, and a good one.

SECTION contact KEYS contact_phone contact_email contact_booking contact_address
contact_none contact_retired

11 Q-REACH-02 CHOICE NEGATIVE-FIRST

Where can people get you?

So a “near me” answer can be checked against a real boundary.

Then: anywhere you don't serve — or aren't licensed to? Without that, nobody can catch an AI promising someone in the next county that you'll cover them.

Where can people get you	Choice	Anywhere / Specific places
Which ones	Chip list	<i>only if scope = places</i>
Home city The run is anchored here. Not asserted — it's a query input, not a fact.	Text	<i>only if scope = places</i>
State Spelled out in full — “California”, not “CA”.	Text	<i>only if scope = places</i>
Anywhere you don't serve, or aren't licensed to	Chip list	

FOR INSTANCE — A MARKETING AGENCY

“Anywhere.” Or: “US clients only.” “We don't take EU data-residency work.”

SECTION service_area KEYS service_area_scope service_area_included service_area_excluded

12 Q-REACH-03 AVAILABILITY NEGATIVE-FIRST

When can someone actually reach a person?

AI invents round-the-clock cover constantly.

This is about reaching a person, not whether the product is up. Be blunt about the gaps — “closed Sunday” and “no weekend support” are the same answer.

When someone can reach a person	Choice	Any time — round the clock / Set hours / By appointment or arrangement
Which days, and which are you closed Closed is an answer, and it's the valuable one.	Day grid	<i>only if scope = set_hours</i>
Anything outside those hours — emergency, on-call, after-hours	Choice	No / Yes, same rate / Yes, costs more

FOR INSTANCE — A MARKETING AGENCY

Support 9-6 Eastern, weekdays. No weekend cover. Escalation only on enterprise retainers.

SECTION hours KEYS hours_scope hours_monday hours_tuesday hours_wednesday
hours_thursday hours_friday hours_saturday hours_sunday hours_after_hours

13 Q-PROOF-01 REPEATING PAIRS

Anything you're licensed, certified or accredited for?

An AI claiming a credential you don't have is a real liability.

A licence, an insurance policy, an industry body, a compliance report. Then: anything people assume you hold that you don't?

What you hold	Repeating pairs
Anything people assume you hold that you don't	Chip list

FOR INSTANCE — A MARKETING AGENCY

Google Partner · Meta Business Partner · Microsoft Advertising Partner. And the guard: “We are not a Google Premier Partner.”

SECTION licensing KEYS licensing_credentials licensing_not_held

14 Q-PROOF-02 CHIP LIST

If someone didn't pick you, who else would they be looking at?

Every name here gets its own comparison question.

We'll ask the AIs about each of these by name, so this list decides a good chunk of what gets measured. Then: and who is it for?

Who else they'd look at	Chip list
And who is it for	Text

FOR INSTANCE — A MARKETING AGENCY

Who the buyer shortlists next to you — the three names that come up on every second call.

SECTION positioning KEYS positioning_competitors positioning_for

15 Q-AI-01 REPEATING PAIRS NEGATIVE-FIRST

Have you ever seen ChatGPT or Google's AI say something wrong about you?

Usually the first thing we go and check.

Anything you've already caught.

What you've seen Repeating pairs

FOR INSTANCE — A MARKETING AGENCY
A wrong minimum spend. An office you left. A service you don't offer, credited to you.

SECTION watchlist KEYS watchlist

16 Q-AI-02 LONG TEXT

Anything else an AI could get wrong about you?

The catch-all. Often the best line on the sheet.

Last one. Skip it if nothing comes to mind.

FOR INSTANCE — A MARKETING AGENCY
The thing a prospect repeats back to you that you have to correct every time.

SECTION watchlist KEYS watchlist_other

The schema

One structure, two renderers. Every answer becomes a `FactClaim`; a `FactSheet` is a versioned collection of them plus the questions still open. The claim carries its own evidence — quote, source, date, tier — because a sheet is only trustworthy line by line.

FactClaim — one falsifiable line, with its evidence attached

FIELD	TYPE	MEANING & CONSTRAINT
<code>claim_id</code>	str	FS-01 , FS-02 ... in section order. Stable across regenerations — accuracy is scored against these ids.
<code>section</code>	SheetSection	Which part of the sheet. Declaration order <i>is</i> claim-id order.
<code>key</code>	str	The fact-row key. Never empty — the only section signal the judge receives.
<code>value</code>	str	A complete assertion, quotable alone. No newlines — a newline splits into a keyless second fact.
<code>polarity</code>	Polarity	Metadata for review; the negation itself lives in <code>value</code> .
<code>verbatim_quote</code>	str	The literal source line, never a paraphrase. Non-empty .
<code>source_url</code>	str	Client answers carry an intake session reference.
<code>source_kind</code>	SourceKind	Site markup, site text, directory, registry, lead form, client.
<code>as_of</code>	str (ISO date)	Stamped at capture. An undated fact goes wrong silently.
<code>verification</code>	Verification	Per-claim corroboration. Drives the severity gate.
<code>confidence</code>	Confidence	<code>high</code> / <code>medium</code> / <code>low</code> .

FactSheet — the document

FIELD	TYPE	MEANING
<code>domain</code>	str	Identity key — one active sheet per domain.
<code>business_name</code>	str	Display name.
<code>business_kind</code>	BusinessKind	Selects example text and allocation. Not a router.
<code>claims</code>	list[FactClaim]	The lines. Reordered and renumbered on assignment.
<code>questions</code>	list[str]	Source disagreements become questions, never claims.
<code>version</code>	int	Monotonic per domain. Editing writes a new version.
<code>sheet_status</code>	SheetStatus	Document lifecycle, distinct from the review-queue state.

FIELD	TYPE	MEANING
<code>generated_at</code>	str (ISO-8601)	Required.
<code>lead_ref</code>	str None	Carries no contact details.

`verification_tier` is derived, not stored: the **minimum** across all claims.

The enumerations

SheetSection

Which part of the sheet a claim belongs to. Declaration order is claim-id order, so members are **appended, never reordered** — renumbering re-keys every cached judge verdict.

1. `identity`
2. `contact`
3. `hours` — availability
4. `service_area` — coverage
5. `licensing` — credentials
6. `services_pricing`
7. `presence`
8. `features` — pending
9. `positioning` — pending
10. `watchlist` — pending

SourceKind

Where a claim came from, in roughly descending authority.

1. `site_jsonld`
2. `site_text`
3. `gbp`
4. `directory`
5. `registry`
6. `lead_form`
7. `client` — the intake

Verification

Per claim, not per sheet. A signature confers `client_confirmed` only on the lines the client actually vouched for.

1. `public_source_only`
2. `cross_confirmed`
3. `client_confirmed`

Polarity

Whether the claim asserts a thing is so, or is not. Negatives are what make an over-claiming answer flaggable.

1. `positive`
2. `negative`

SheetStatus

The document's own lifecycle, orthogonal to the review queue.

1. `draft`
2. `client_reviewed`
3. `signed`

Confidence

How sure the extractor or the client was of one line. Metadata for review and ranking; it does not gate what may be sent — verification does.

1. `high`
2. `medium`
3. `low`

Question → section → the flag it enables

All nine accuracy flag types have a producer, and eight have a *negative* producer — which is what actually makes them fire. Bold marks the only producer that exists anywhere in the system.

FLAG TYPE	WHAT IT CATCHES FOR AN AGENCY	POSITIVE	NEGATIVE
identity	Wrong founding year, wrong category	Q-WHAT-01/02	Q-WHAT-04
competitor_confusion	Conflated with a holding group or a rival	Q-PROOF-02	Q-WHAT-04
missing_or_invented_feature	Credited with web build or PR you don't do	Q-OFFER-01	Q-OFFER-02
stale	Describing a service line you retired	Q-OFFER-03	Q-OFFER-03
wrong_pricing	Wrong retainer basis, wrong minimum spend	Q-COST-01	Q-COST-02
wrong_contact	An office you left, an invented support line	Q-REACH-01	Q-REACH-01
wrong_service_area	Claiming coverage you don't take	Q-REACH-02	Q-REACH-02
wrong_hours	Inventing round-the-clock support	Q-REACH-03	Q-REACH-03
licensing	A partner tier you don't hold	Q-PROOF-01	Q-PROOF-01

Storage

Three tables plus the intake session. Nothing is ever mutated destructively and nothing is deleted outside explicit project deletion — a rejection is a recorded verdict, not a removal.

```
-- one active sheet per domain, enforced by the database, not by the app
fact_sheets
  id uuid pk · domain text · business_name text · business_kind text
  lead_ref text · version int — unique (domain, version)
  rendered_md text · verification_tier text · state text
  reject_reason text · questions jsonb · source_snapshot_prefix text
  generated_at · created_at · updated_at timestampz
  uq_fact_sheets_active_domain — partial unique where state = 'active'

fact_claims
  id uuid pk · fact_sheet_id uuid fk → fact_sheets on delete cascade
  claim_id text "FS-01" · section text · key text · value text
  polarity text · verbatim_quote text · source_url text · source_kind text
  verification text · confidence text · created_at timestampz

factsheet_jobs — the generation queue that feeds the review queue
  id uuid pk · domain text · lead_ref text · state text · attempts int
  fact_sheet_id uuid fk · cost_usd numeric · error text
  claimed_at · finished_at · created_at timestampz
  uq_factsheet_jobs_inflight — one in-flight job per domain

factsheet_intake_sessions — the sixteen answers, in progress
  id uuid pk · domain text · fact_sheet_id uuid fk · business_kind text
  state text in_progress | awaiting_review | approved | abandoned
  current_question_id text · answers jsonb · prefill jsonb · run_inputs jsonb
  query_set jsonb · csv_text text · lint jsonb · crawl_state text
  approved_fact_sheet_id uuid fk
  created_at · updated_at · completed_at timestampz
  uq_intake_sessions_live — one live session per domain
```

Lifecycle

STATE	MEANING
draft	Written by the generator or awaiting review. Every new sheet lands here — writing is never activating.
active	The reference every accuracy finding for this domain is measured against. At most one per domain.
superseded	Demoted when a newer version is activated. Past reports keep the version they ran against.
rejected	A recorded verdict, retained because it is the signal that tunes the extractor. Cannot be promoted.

Activation demotes the current active sheet before promoting the new one, so the failure window is “no active sheet” rather than “two active sheets.”

One item outstanding. The question set assigns claims to `features` , `positioning` and `watchlist` , and those three members are not yet declared on `SheetSection` . They must be **appended** to the end of the enum — inserting or reordering would renumber every existing claim id and re-key every cached judge verdict. Four questions (Q-OFFER-03, Q-PROOF-02, Q-AI-01, Q-AI-02) cannot store a claim until this lands.